

Deploying Stateless Application with Deployment Objects.

Kubectl - kubectl controls the Kubernetes cluster manager

You can run an application by creating a Kubernetes Deployment object, and you can describe a Deployment in a YAML file. For example, this YAML file describes a Deployment that runs the **<your-docker-hub-username>/dell:v1** Docker image

Replace **<your-docker-hub-username>/dell:v1** with the image name, repository name and the tag from your **docker hub** account that you pushed to docker hub in **docker labs**

1. SSH to your AWS Instance and make a dir /home/<your-user-name>/application

```
$ cd /home/devops
$ sudo mkdir application
$ cd application/
```

Create a New Deployment file starting with your name.

1. \$ curl -k <https://pastebin.com/raw/krKi6xsh> > <your-name>-deployment.yaml

```
$ vim <your-name>-deployment.yaml
```

#Note : press 'i' to start the edit mode in the vim editor. Replace <your-name> with your name and then save and exit by (:wq (enter))

2. Create a Deployment based on the YAML file:

```
$ kubectl apply -f <your-name>-deployment.yaml
```

3. Display information about the Deployment:

```
$ kubectl describe deployment <your-name>-deployment
```

The output is similar to this:

```
Name:      <your-name>-deployment
Namespace:  default
CreationTimestamp: Tue, 30 Aug 2016 18:11:37 -0700
Labels:     app=dell
Annotations: deployment.kubernetes.io/revision=1
Selector:   app=dell
Replicas:   2 desired | 2 updated | 2 total | 2 available | 0 unavailable
StrategyType: RollingUpdate
```

MinReadySeconds: 0
RollingUpdateStrategy: 1 max unavailable, 1 max surge
Pod Template:
Labels: app=dell
Containers:
nginx:
Image: asyed755/dell:v1
Port: 80/TCP
Environment: <none>
Mounts: <none>
Volumes: <none>
Conditions:
Type Status Reason

Available True MinimumReplicasAvailable
Progressing True NewReplicaSetAvailable
OldReplicaSets: <none>
NewReplicaSet: nginx-deployment-1771418926 (2/2 replicas created)
No events.

4. List the pods created by the deployment:

\$ kubectl get pods -l app=<your-app-name>

The output is similar to this:

NAME	READY	STATUS	RESTARTS	AGE
dell-deployment-1471416983-7o5ac	1/1	Running	0	16h
dell-deployment-1541148254-318ad	1/1	Running	0	16h

5. To display information about a pod:

\$ kubectl describe pod <pod-name>

6. Expose the Deployment with the below command.

\$ kubectl expose deployment <your-app-name> --type=NodePort --name=<your-service-name>

7. Check the Node details where the POD has been deployed

\$ kubectl get pod -o wide | grep <your-app-name>

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
arshad-5fdff48b48-7r4pg	1/1	Running	0	2m	100.96.4.43	ip-172-20-55-125.ec2.internal
arshad-5fdff48b48-gg7j4	1/1	Running	0	2m	100.96.4.44	ip-172-20-55-125.ec2.internal
rajni-deployment-d667	1/1	Running	0	3h	100.96.4.13	ip-172-20-55-125.ec2.internal

ip-172-20-55-125.ec2.internal is the Worker node where the app has been deployed

8. Login to the **AWS Console** and Copy the **Public IP** of the NODE where the POD has been deployed to access the application.

Access the application as shown below

<http://<node-public-ip>:NodePort> ##see step 5 to check the port

Example

<http://18.208.206.161:32587/s>

